

PARTH KHARE

✉ khare.parth2112@gmail.com [in LinkedIn](#) [📁 Portfolio](#)

Summary

AI Engineer and full-stack developer building production LLM-integrated systems — from retrieval and verification pipelines to REST APIs and deployment. Comfortable owning the full stack: LLM orchestration, backend architecture, database design, and containerized delivery. Shipped multi-tenant SaaS with authentication, payments, and real-time frontend sync; deployed ML systems with drift monitoring in production. Looking to bring that same ownership to an AI Engineering or Full-stack role.

Skills

Languages: Python, SQL, JavaScript, TypeScript, HTML/CSS

Libraries & Framework: React, Flask, Scikit-learn, Streamlit

Backend & APIs: Fast API, REST APIs, JWT/RBAC, SQLAlchemy, Webhooks(HMAC)

Database: PostgreSQL, Redis, MySQL, MongoDB

DevOps & Tools: Docker, Git, AWS, Razorpay API, Power BI

Projects

Veritas AI – Anti-Hallucination & AI Output Verification API | [View Project](#)  **In Development**

Python, FastAPI, sentence-transformer, PostgreSQL, Redis, Docker

- Building a model-agnostic REST API that verifies LLM outputs through a layered pipeline — syntactic contradiction/citation checks, NLI-based faithfulness scoring against source context, and LLM-driven atomic claim decomposition — returning a structured trust-score and per-claim verdict.
- Designing the system to be pluggable across domains (general/code/medical/legal) and model-agnostic by construction, so it can sit between any LLM (GPT, Claude, Gemini, Llama) and end users as a drop-in verification layer.

AI Resume Screener (B2B2C SaaS Application) | [View Project](#)  **April 2026 – June 2026**

React, TypeScript, FastAPI, PostgreSQL, Docker, Gemini API

- Built a **multi-tenant AI SaaS platform** automating resume screening - Gemini API + PDF parsing generates match scores, skill gaps, and structured candidate insights.
- Architected secure backend with **FastAPI, PostgreSQL, RBAC, and JWT auth**; integrated Razorpay payments with **HMAC webhook** verification and real-time frontend state sync.

Customer Churn Prediction System | [View project](#)  **Jan 2026 – April 2026**

Python, FastAPI, Streamlit, XGBoost, Docker

- Engineered end-to-end churn prediction system on 7,000+ telecom customers; achieved ROC-AUC 0.84 and 88% recall, identifying 9 in 10 at-risk customers before they churn.
- Reduced false negatives by 23% over baseline logistic regression by switching to threshold-optimized XGBoost (threshold=0.22 via cost matrix: FN=10×FP), directly minimizing missed-churn business cost.
- Deployed containerized REST API (FastAPI + Docker) serving real-time predictions with Redis caching, cutting redundant inference to zero for repeat queries.
- Built Streamlit analytics dashboard with SHAP waterfall explanations, enabling non-technical stakeholders to identify top churn drivers per customer without data team involvement.
- Implemented input validation, structured JSON logging, and PSI-based drift monitoring across 3 key features, reducing silent model degradation risk in production.

Education

B. Tech. in Computer Science and Engineering

May 2026

Vellore Institute of Technology, Bhopal

CGPA: 7.76 / 10

Certifications & Training

- Oracle Cloud Infrastructure AI Foundation Associate Certificate.
- Complete Data Analytics Bootcamp – Udemy
- Deloitte Australia Data Analytics Job Simulation Certificate
- AWS Solution Architect Training – Ethnus